

Making an FPGA-based Digital Down-Converter: Ay121 Lab Instructions

March 8, 2011

For the next two weeks, we'll be exploring real-time digital signal processing with state-of-the-art digital computing hardware. We're going to use a ROACH board designed by the Collaboration for Astronomy Signal Processing and Electronics Research (CASPER). This board employs a massively parallel Field-Programmable Gate Array processor, and could, with the exception of a front-end amplifier, possibly replace the entire wall of hardware in the radio lab.

In this lab, we're going to replace only a part of that hardware; we'll turn the ROACH into a digital mixer and filter (which we together will call a down-converter). Each *group* will work together on programming the ROACH board using the Simulink programming language. Once a design is compiled, each group will setup a noise test on the bench with the ROACH running their compiled design. However, each *individual* will choose her own coefficients to use in the FIR filter we implement. Each individual will then record her own filtered and unfiltered data, and use this data to measure the per-frequency response of her filter coefficients. The individual will compare this measured response to the analytic response she predicts through analysis in IDL.

1 Accessing Simulink

1.1 Setting up a VNC Session

From any *nix computer:

```
$ ssh username@host.computer.edu
    (enter password)
$ vncserver
    (write down XX of "desktop is host.computer.edu:XX")
$ exit
```

This will set up a VNC session for you to log into in the future. Remember your session number (XX).

When you are done with your session, it would be neighborly of you to free up the computer resources associated with your vnc session:

```
$ ssh username@host.computer.edu
$ vncserver -kill :XX
    (where XX is your VNC session from above)
$ exit
```

Thanks!

1.2 Accessing Your VNC Session

From any *nix computer:

```
$ vncviewer -via username@host.computer.edu localhost:XX
    (ssh password)
    (vnc password)
```

All future commands are to be typed inside your VNC session.

1.3 About Simulink on Linux

The machine you are logged into is running a Linux desktop. It behaves somewhat differently than other desktop environments you may be used to. Because Simulink sometimes opens windows that are larger than the size of the desktop, it may be useful to know that you can move these windows around by right-clicking, selecting "Move", and then dragging the window around.

1.4 Starting Matlab/Simulink

Make sure there is a file "startup.m" in your directory. If it does not exist, copy it there from /opt/casper/startup.m. This file is required to automatically load some signal processing libraries when you start up the design tools. To start Matlab/Simulink, type:

```
$ simmy
```

This program is a script that sets up the licenses for all the tools we'll be using, and then executes the program "matlab".

2 Introduction to Simulink and ROACH

We have a tutorial at <http://casper.berkeley.edu/astrobaki/index.php/CasperTutorial01> that will run you through the basics of FPGA design with Simulink. Skip section 8 of the tutorial—it runs through some software (KATCP) that you won't need use.

A few details from that tutorial are different for the setup in our lab. For one, the ROACH we use has an "lx110t" FPGA, which is different from the "sx95t" that is selected in the *XSG Core Config* block in the tutorial. Our ROACH in the lab is at IP address 10.32.92.113, but can also be called "roach". Whenever you get to the part about sending your *.bof file over to the ROACH, you'll need to first scp the file to your account in the lab, and then copy the file from there to root@roach:/boffiles. To read/write from the registers in your design, you can ssh in from any computer in the lab, but for an authentic experience, you might want to ssh in from a computer on the bench. That way, you can see the blinken light, and you won't have the problem of two people trying to program the FPGA at the same time.

Here are a couple other pointers that weren't in the tutorial:

- In the window containing your Simulink design, choose *Format* → *Port/Signal Displays* → *Port Data Types*. Although this will complain if your design isn't complete (i.e. if there are errors), it will tell you the data type of each wire of a working design. This can really help you follow how your design works and catch bugs before compiling. If you change your design, hit Ctrl-D to update the labels.

3 Capturing ADC Samples

Now that you are more-or-less familiar with the Simulink programming language, it's time to create a design that captures raw data samples from the Analog-to-Digital Converter attached to the ROACH.

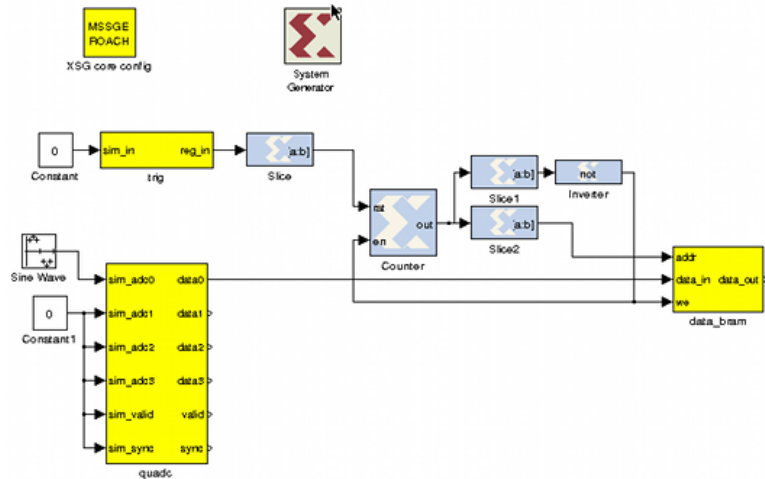
3.1 The Simulink Design

The ADC we are using is a Quad-ADC or "quadc", so named because it actually digitizes 4 independent signals. We'll only be using one signal. To interface to the QADC, drag in a **quadc** yellow block from the BEE_XPS System Blockset. Providing appropriate simulation inputs, connect the "data0" output line to a **Shared BRAM**, and rename that BRAM to "data_bram". Configure the Block RAM (BRAM) to hold a lot (at least 2048) data samples.

Next, you'll want to create a mechanism for triggering a capture of data into the Shared BRAM. This circuit should walk through all of the addresses in the BRAM sequentially, writing a new ADC sample into each slot. Once the last address is reached, the counter should halt, and data should stop being written. The system should then wait until you reconfigure a **Software Register** (renamed to "trig") to trigger a new data capture.

A picture of a working design has been included as a reference should you get stuck. However, you'll have to figure out the block parameters for yourself. One thing to be sure of: make sure the **XSG Core Config** block in your design is set for "User IP Clock Source" = "adc0_clk". This uses the ADC sample clock to clock the FPGA, so everything marches through the FPGA at the rate that samples arrive.

Remember to simulate your design to ensure everything is working before you compile!



3.2 Running the Design

After compiling, transfer the *.bof file over to the ROACH as before, and run it. Make sure the sample clock to the ADC is set to 200 MHz, and is on before you program the FPGA. FPGAs behave very erratically if they aren't clocked properly; if you program the FPGA and then turn on the clock, you may see gibberish, or even nearly-correct behavior that occasionally goes whacky. Also check that you have attached a noise signal into the ADC, so that you get something interesting in your BRAM. You can now write appropriate values to "trig" to initiate a data capture. You can use `hd` to look at values in "data_bram" and make sure things make sense. When you are happy with the contents of the BRAM, use your account at the lab to scp the contents of the BRAM to a computer with IDL:

```
$ scp root@roach:/proc/PID/hw/ioreg/data_bram .
```

Create a power spectrum of the noise you recorded, with a number of channels at least a factor of 8 smaller than the number of ADC samples you captured. Use the fact that you can compute several such power spectra over the window of samples you acquired to beat down the noise in your measurement of the noise power spectrum.

1. Plot the power spectrum of the noise you recorded, integrated over multiple DFT windows

You might be noticing a huge spike of power in the frequency 0 bin. This comes from a voltage bias at the input of the ADC. The specification for this ADC allows for up to 3 bits of signal bias. You'll probably want to zero out this bin to make your plots look better. Remember to omit this bin from your coming analysis.

4 Applying a Summing Filter

Modify (but save to a new file) your design to capture ADC samples, so that two adjacent input time samples are summed to produce each output sample. States more precisely, produce the output S such that:

$$S_i = X_{i-1} + X_{i-2} \quad (1)$$

where X_i are input samples. The latency of your summing filter is unimportant.

Before programming the ROACH with this design:

1. Can you predict a priori what frequency response this time-domain summing filter will produce (hint: it *is* an analytic function)?
2. If you don't know yet what the response function will be, simulate it with noise in IDL to obtain a low-noise predicted filter response (in units of power). If you do know the response, generate the exact filter response.

Follow the same steps you used for the ADC capture design, inject the same noise source and capture a new set of time domain samples drawn from after the application of the summing filter.

1. Plot the power spectrum of the filtered noise you recorded, integrated over multiple DFT windows

You now have two sets of time-domain samples, each drawn from the same noise source and digitized with the same ADC, but one has had a digital filter applied. Given that the noise source and the ADC do not have a perfectly flat response versus frequency:

1. Use your measured power spectra to isolate the inherent response of your summing filter
2. Overlay on your measured filter response the filter response you predicted
3. Is your prediction accurate to within the noise of your measurements? What are your error bars?

5 Create an FIR Filter

Implement an 8-tap FIR filter:

$$S_i = \sum_{n=0}^7 C_n X_{i-n} \quad (2)$$

The coefficients of this filter, $C(t)$, will be a function that convolves the input signal, $X(t)$, to produce the time-domain filtered response $S(t)$:

$$S(t) = C(t) * X(t) \quad (3)$$

Since a convolution in time-domain is multiplication in the frequency domain, this convolution is equivalent to:

$$\hat{S}(\omega) = \hat{C}(\omega) \cdot \hat{X}(\omega) \quad (4)$$

You will want to implement the FIR coefficients with run-time programmable with **Software Registers**. You may find it useful in Simulink to highlight a set of blocks and right-click → *Create Subsystem*. This will create a block that you can rename, copy, and paste just as any other block, but underneath contains your blocks. You can open it by double-clicking. A stylistic point: I recommend against putting yellow blocks in subsystems (although I do break this rule sometimes). The rationale is that once you compile your design, you'll want to know at a glance what the interfaces are. If you hide your interfaces in subsystems, it's hard for a user to see quickly how they are supposed to interact with the design.

Copy your previous design for capturing time-domain samples. Before programming the ROACH with this design:

1. Choose coefficients that implement a 5/8-band filter (a filter that, if you divided the band into 8 complex channels, would extract the 5 channels centered around 0):

Choose the frequency-domain response you want for your filter (i.e. the function you would like to multiply the frequency spectrum of your noise by to "filter" it)

Compute the real-valued, time-domain coefficients that implement that filter, keeping in mind that:

Multiplication in the frequency domain is a convolution in the time domain

The Fourier transform of a real-valued signal has the property that: $\hat{f}(\omega) = \hat{f}^*(-\omega)$

You will want to think carefully about where the 0 frequency bin is

The FFT puts negative frequencies after positive frequencies in your array. Similarly, when you take the inverse, it will put negative times after positive times. When implementing your coefficients in an FIR, though, negative times *have* to operate on samples that arrive before positive times.

2. In what order are you going to write these coefficient into the software registers of your FIR filter?
3. Plot the predicted filter response for your coefficients at a frequency resolution much finer than an 8-channel DFT produces. The best way to do this will be to add more time-domain samples to your coefficients. We don't want to introduce any new signal, though, so just add zeros to pad your coefficients out to 64 samples, and then transform the result back into the time domain.
4. Why aren't the passband and stopband flat?

Now collect data from the noise source, filtered through the FIR filter programmed with your coefficients. As before:

1. Isolate the inherent response of your FIR filter

2. Compare your measured filter response to filter response you predicted
3. Is there a way you can think of to create the exact same FIR filter (which has symmetric positive and negative frequency response) using half the number of multipliers and software registers? There's no need to actually do it, unless you want to.

6 Add a Mixer

As a final touch, we'll add a digital mixer. Compared to an FIR filter, this will be pretty quick. The only trick is creating a digital sine and cosine. We'll take the easy way out here, and use the **SineCosine** block from the *Xilinx Blockset* → *Math* library. The "theta" input should just be a counter. The number of bits in the counter will determine the period of the sine/cosine wave generated. Although it looks fancy, it's really just a BRAM loaded with coefficients, and "theta" is the address. For a reason I'm not privy to, the "theta" port must be at least 3 bits. If you want something that cycles faster than a 3-bit counter, you'll have to concatenate a counter and zeros to pad the signal out to 3 bits.

1. If you will be clocking this design at 200 MHz, and you want to mix at 50 MHz, what should the sample period of your sine/cosine be?

Mixing with a sine/cosine gives you In-phase and Quadrature signals, which we'll take to be a complex number. Our FIR filter was only equipped to handle real numbers. Fortunately, extending the FIR filter for complex numbers should really just be a matter of copy-paste, especially if you kept your coefficient registers on the top level. We would like to keep just one set of coefficients to avoid unnecessary interfaces. You'll also notice that we need to rework the **Shared BRAM** block, since we now have real and imaginary components to record. You can either:

- cop out and add a second bram that attaches to the same address and write-enable so that it records data at the *exact* same time as the other one, or
- keep your design slim by trimming and concatenating the signal to fit into the 32-bit data_in port. If you go this route, pay attention to your data types and don't forget to make your Shared BRAM of unsigned data type.

A major challenge in creating an analog down-converter lies in creating sine/cosine waves that are perfectly in phase with one another, and in carefully matching the filters on the real and imaginary components of the mixer product. Any mismatch in phase or response creates an "mirroring" problem, whereby a positive frequency will have a small mirror image at the corresponding negative frequency. Perfectly in-phase sine/cosine signals and perfectly matched filters come for free in our digital down-converter. This is nothing to sneeze at. Digital down-converters have only recently become viable at radio frequencies, and they are fast supplanting their analog forebearers.

Finally, let's prove the down converter works. Create a new signal input that is the sum of noise and a 60 MHz sine wave. Record the down-converted result, and:

1. Plot the time-averaged power spectrum of the signal.
2. At what frequency should the output tone be? How many samples the period of this tone?
3. What is the range (in MHz) of the passband of your FIR filter?
4. Through your FIR filter, the number of bits increased with each operation. Why was it okay, for those who fit the complex data into one Shared BRAM output, to reduce the number of bits? What would be the minimum number of bits you need to output the signal without significant loss in precision?
5. If you knew that the signal you were processing was noise-dominated, and your FIR filter's coefficients were set to quarter the bandwidth, could you shave any more off the minimum number of bits needed in your complex number? Explain.

7 Conclusion

Hopefully you've gotten a taste for how digital signal processing works on an FPGA. Simulink is a fun graphical programming tool that captures the inherently parallel nature of circuit design. Programming for parallel processing is a lot different than the iterative processing model we are used to thinking about when programming CPUs. If you enjoy the challenge of programming FPGAs and are thirsty for more, the CASPER group supports undergraduate research in the field of radio astronomy instrumentation and digital signal processing. Many of us use FPGA programming and instrument design as tools for furthering our science objectives in radio astronomy, and welcome interested students. Feel free to contact Aaron Parsons (aparsons@astron) for more information.